


Loop engineering on one page

From Joe Black's talk, She Loves Data Vietnam, 23 September 2026

Prompt engineering is one instruction you babysit. Loop engineering is a job you hand over: you write down what good looks like, and a small team of AI agents does the work, checks it against that, fixes what failed, and only then comes back to you.



THE FIVE PARTS (NUMBERED TO MATCH THE DRAWING)

- 1 What sets it off.** You handing it the job. Or a schedule, a form, a file landing, if the job repeats.
- 2 What rules it reads.** Your "what good looks like", read every lap. What not to touch.
- 3 Who checks the work.** A different agent from the one that did it. Never the doer.
- 4 How it knows it's finished.** A written test the machine can run: counts, links, numbers sourced, tests green.
- 5 Where it keeps its notes.** A document it reads first and updates last, so a crash never means starting over.

Before anyone leaves the room: the goal and what done looks like; what sets it off; the rules; who does it and who checks it; the sign-off test; where the notes live; and where it must stop and ask you.

To loop or not to loop: Would it eat your day going back and forth? Can you write down what good looks like? Will it pay back the half hour it takes to set up? Can it see its own results? Four yeses, build it. One no, it's a prompt.

One line to remember: A scheduled prompt is not a loop. A scheduled loop is.

The loop you build tonight

Paste this into ChatGPT, Gemini or Claude, inside a project. Change the field line. Run it once. Happy? Put it on a schedule.

Do this job. Run it every Monday at 7:30am.

JOB: Find what changed in my field in the last 7 days that I should know about.

My field: [write yours here, e.g. HR and people management for companies in Vietnam]

My situation: [one line, e.g. 200 people, one office, no foreign staff]

RULES (what good looks like):

- Only things first published inside the last 7 days, with a link I can open.
If an article restates an older report, the date is the report's date.
- 3 to 5 items, ranked by how much they change what I do this week.
If fewer than 3 real items exist, the report says "Quiet week" and nothing else. Never pad.
- Each item: what it is (2 lines max), why it matters to me (1 line, concrete), the date, the link.
- Skip anything in the ALREADY SEEN file.

CHECK THE WORK, as a different agent: a skeptical editor who did not write the list.

For each item confirm: (1) first published inside the last 7 days, (2) the link is real and points to that item, (3) the "why it matters" line names something concrete, not "this could be interesting". Anything that fails goes back to the writer to replace or drop. Keep going until every item passes.

DONE WHEN: every item passes the editor's three checks, and the ALREADY SEEN file has this week's items added.

NOTES: keep a file called ALREADY SEEN in this project, one line per item: title, one-line note, date. Read it before you start. Add this week's items when you finish. Anything dropped, and why, goes in the notes, not in the report.

REPORT: one clean page. The date at the top, the items, nothing else. No working shown.

The five parts, marked: JOB is the goal. The first line is what sets it off. RULES is what it reads. CHECK THE WORK is the checker, a second agent. DONE WHEN is the gate. NOTES is where it keeps its notes. REPORT is what you see.

Where to put it: ChatGPT: a Project holding the rules, then Tasks for the schedule. Gemini: a Gem, then a scheduled action. Claude: a Project, then a scheduled task in the desktop app. If your app can't write to a file yet, it prints the seen list and you paste it into the file. Training wheels; the loop still works.

Three loops to copy: the article that passes three reviewer agents before you see it; the report where every number must trace to its source or it goes back; the monthly check of your website against the best out there.