

Loop Engineering

Stop asking AI questions. Start handing
it work.

Joe Black · She Loves Data Vietnam · 23 September 2026



**Hands up if you
used AI today.**

Keep them up if you still did the work.

A question gets
you an answer.

A job gets you a result.

Prompt engineering: one instruction you babysit.

Loop engineering: a job you hand over.



"I don't prompt Claude anymore. There are loops running, and it's the loops that prompt Claude. My job is to write the loops."

Boris Cherny, who built Claude Code, June 2026

So what the hell is an agent?

A chat is a conversation.

An agent is a colleague. A job title, its own instructions, its own tools, and a job it goes off and does.

Same AI underneath. Different job description.

No coffee breaks. No opinions about the aircon.

You already know how to hand over a job.

Picture the meeting where you brief a team. Nobody leaves the room until these are settled.

The goal

What are we making, and what does done look like.

What sets it off

You handing it the job. Or a schedule, if the job repeats.

What rules it reads

What good looks like. What not to touch.

Who does the work, and who checks it

Never the same person.

How it knows it's finished


The sign-off test, written down.

Where it keeps its notes

So nobody starts over.

That's a loop. You've run them with people for years.


How it runs.



You build it once. Then you step back.

Meetings that end on time.

BufferFlow. A timer I wanted and couldn't buy, so I made one.



Version one was vibe-coded.

It looked like software but...

The sign-in was fake. Nothing saved. The AI button returned the same agenda every time.

The night of 28 August

7:11pm

End of my workday. I wrote what I wanted, in English.


A friend's four fixes. What done looked like. What not to touch. Not a line of code.

Then I got into it.

I went to bed.

1:30am. Best management decision I made all month.

That night, on the loop.



7:11pm started · 1:27am last change · 10:26am merged and live

10:26am. I read what it did, merged it, live.

18 changes shipped.

42 files, 6,400 lines.


Tests: 34 to 213.

A separate reviewer found 13 problems. 12 fixed, 1 argued down.

My longest run since: 11 hours 38 minutes, 27 changes, on a bigger piece of software. They keep getting longer.

BufferFlow repo, git log, 28 to 29 Aug 2026. Longest run: git log, 18 Aug 2026.

The same five in the app you already have



	THE PART	IN CHATGPT, GEMINI OR CLAUDE
1	What sets it off	You handing it the job. Or a schedule, if the job repeats.
2	What rules it reads	A project or custom assistant holding your "what good looks like".
3	Who checks the work	A second pass by a different persona, a second chat, or a second model.
4	How it knows it's finished	A pass or fail checklist it has to print, line by line.
5	Where it keeps its notes	A document it reads first and updates last.

ChatGPT: Projects and Tasks. Gemini: Gems and Scheduled actions. Claude: Projects and Scheduled tasks. A scheduled prompt is not a loop. A scheduled loop is.

Never let the doer grade its own work.

It gives itself an A+. Every time. So would you.

Checking its own work, an AI caught 1 error in 4.


Checking twice, it caught fewer. A fresh reviewer that never saw the work caught the most.

Same AI, fresh eyes, its own instructions: that's a different colleague. A different AI catches more again.

First separate review of BufferFlow, 3 July: 30 problems in code that already worked.

The loop on your handout, first test run: its checker binned a "stat" every article was repeating. The source was a year old.

Song, Cross-Context Review, arXiv, Mar 2026 (150 planted errors, 30 documents) · Xiang et al., arXiv, Jul 2026 · BufferFlow repo, 3 Jul 2026 · handout loop test run, 15 Sep 2026.



Slop compounds.

"One day I actually read it closely. It had been off for weeks."

Prompt engineering: you checked every step.

Loop engineering: you write down what good looks like, you say where it must stop and ask, and you still check the work.

To loop or not to loop. That is the question.

Would it eat your day going back and forth?

Can you write down what good looks like?

Will it pay back the half hour it takes to set up?

Can it see its own results?

Loop it: the report where every number must trace to a source. The article that has to pass three reviewers. Anything you'd otherwise babysit all day.

Don't: one email. One summary. Your wedding speech. Anything you couldn't grade yourself.

Four yeses, build the loop. One no, it's a prompt. Go do it.

Loops worth building

The article

Written, then three reviewer agents: fact-checker, relevance, house style. Round it goes until all three pass. You read one draft.

The report

Every number must trace to its source or it goes back to the writer. You get a page you can sign.

My website

Once a month it reads the best sites out there, checks its ideas against what I've already done or rejected, and hands me a shortlist.

Put a loop on a schedule if the job repeats. Then it's a scheduled loop, not a scheduled prompt, and the output shows it.

The recipe, in five parts

What sets it off	Paste it in and say go. Every Monday at 7:30am, if you put it on a schedule.
What rules it reads	Dated inside the week. A link I can open. 3 to 5 items. Never pad. Skip the seen list.
Who checks the work	A second agent, the skeptical editor, checks every item. No real date or link, it goes back.
How it knows it's finished	Done when every item passes the editor.
Where it keeps its notes	The ALREADY SEEN file. Read first, added to last.

You get one clean page. The working stays in the notes. Full text on the handout.

The handout

Two pages. The theory on one, tonight's recipe on the other.

Scan the code on the last slide.

Loop engineering on one page

From Joe Black's talk, She Loves Data Vietnam, 23 September 2026

Prompt engineering is one instruction you baby! Loop engineering is a job you hand over: you write down what good looks like, and a small system does the work, checks it against that, how what failed, and only then comes back to you.



Yes, done. It comes back to you.

THE FIVE PARTS (NUMBERED TO MATCH THE DRAWING)

- 1 **What sets it off.** A schedule, a request, a feeling.
- 2 **What rules it needs.** Your "what good looks like", read every lap. What not to touch.
- 3 **Who checks the work.** A different agent from the one that d-d-it. Never the deer.
- 4 **How it knows it's finished.** A written test the machine can run. Counts, lists, numbers sourced, tests green.
- 5 **Where it happens to notes.** A document it reads test and updates last, so a crash never means starting over.

Before anyone leaves the room, the goal and what does look like, what sets it off, the rules, who does it and who checks it, the sign-off test, where the notes live, and when it must stop and say you.

The one-minute test, before you build a loop: Do you do it again and again? Could a machine tell if it's right? Is the result worth what it costs to run? Can't see why it failed? Four years, build it. One no, go do it yourself.

joeblackdreamscope.wiki | linkedin.com/in/joeblack

dreamscope

The loop you build tonight

Paste this into a scheduled task in ChatGPT, Gemini or Claude. Change the field line. Ten minutes.

Every Monday at 7:00, do this job.

GO: Find what changed in my field for the last 7 days that I should know about.
my field: Service plans team, e.g., we and people management for companies in various
by location: [see link, e.g., see people, one office, no foreign stuff]

Write your good looks like:

- Only today. First published inside the last 7 days, with a link, I can see.
- It is an article or update or other report, the date is the report's date.
- I do 3 items, ranked by how much they change what I do this week.
- If they don't fit 3, read them all, write "look what we did". Read last.
- Last item what it is. (I know what, why it matters to me (I know, connected, the date, the link, the report, or my last time star).

DO: Now, do the work, as a different person, a checklist reader who can tell what the last.
for each time, write: (1) The date is inside the last 7 days. (2) The link is real and points to that item. (3) The "why it matters" line comes, reading content, with "this is not so interesting", then repeating that rule. (4) The rule you changed and why, one line each.

OK, back: every morning I see how a date and a working link, and one line printed the content is inside the last 7 days with my notes? Then write, ready for me to paste back into this task.

READY: YES (one per line: YES/NO, one line note, date).
- Check and

The five parts, marked: Line 1 sets it off, RULES is what it needs. CHECK YOUR OWN WORK is the checker (stronger, paste the output into a second chat with only that paragraph). DONE WHEN is the gate. ALWAYS SEEN is the note.

Where to put it: ChatGPT, Gemini, make a Project that holds your rules. Define a scheduled action, with a Gem holding the rules. Claude, a scheduled task in the desktop app, with a Project holding the rules. New rules change the five parts don't.

Three small loops to copy: Every Monday, what changed in my field. Every month, what's changed on my website against the best out there. Every Friday, the update for my boss, draft it, check every number against the sheet, fix, hand me the field.

One line to remember: A schedule is just what sets it off. The check is what makes it a loop.

joeblackdreamscope.wiki | linkedin.com/in/joeblack

dreamscope

The key is knowing what good looks like.

Stop asking. Start handing it work.

Slides and the handout: scan the code. Joe Black · joe@dreamscope.win

